

Raju Nalla

 (+91) 9100173325 |  nallaraju00000@gmail.com |  [raju-nalla](#) |  [raju-nalla](#) |  [raju-nalla.github.io](#)

Professional Summary

Azure & Snowflake Data Engineer with **4+ years of experience** designing scalable data platforms, cloud **ETL/ELT** pipelines, and modern analytics solutions. Skilled in **Azure Data Factory, Azure Databricks, Snowflake, PySpark, SQL**, cloud migrations, and enterprise data engineering.

Experienced in building **metadata-driven data pipelines**, implementing **Medallion Architecture, Delta Lake**, and **incremental data processing** to deliver high-performance, reliable data solutions. Strong expertise in **SQL Server to Snowflake migrations, performance optimization, data quality validation**, and collaborating with cross-functional teams to deliver business-driven analytics on Azure cloud.

Education

Bachelor of Technology, Civil Engineering
Jawaharlal Nehru Technological University, Hyderabad, TS, India

Aug 2018 - Dec 2021
GPA: 7.25/10.0

Technical Skills

Programming Languages: Python, SQL, PySpark, Spark SQL, T-SQL.

Cloud & Data Engineering: Microsoft Azure (Azure Data Factory, Azure Databricks, Azure Synapse Analytics, ADLS Gen2, Azure SQL Database, Azure Key Vault), Snowflake (Snowpipe, Snowpark, Streams, Tasks, Time Travel, Virtual Warehouses), AWS (S3, Glue, Lambda), Google Cloud Platform (BigQuery, Cloud Storage).

Data Engineering & Big Data: Apache Spark, Delta Lake, ETL/ELT Pipelines, Data Warehousing, Medallion Architecture, Incremental & CDC Processing, SCD Type 1 & Type 2, Data Modeling, Performance Tuning.

Databases & Analytics: SQL Server, Snowflake, Azure SQL Database, Oracle, PostgreSQL, Power BI, Tableau.

DevOps & Collaboration: Git, GitHub, Azure DevOps, CI/CD Pipelines, YAML, Databricks Workflows, Terraform, SnowConvert AI.

Certificates

- **Databricks Certified Data Engineer Associate** | Mar 2026 |
- **Microsoft Certified: Fabric Data Engineer Associate (DP-700)** | May 2026 |
- **SnowPro Core Certification** | Jul 2026 |

Professional Experience

Azure Data Engineer
Manufacturing Client

May 2024 - Present
Hyderabad, India.

- Designed and developed **Azure Data Factory (ADF)** pipelines to ingest data from **SQL Server, REST APIs**, and **Flat Files** into **Azure Data Lake Storage Gen2 (ADLS Gen2)** for enterprise analytics.
- Engineered scalable **ETL/ELT** pipelines using **Azure Databricks, PySpark, Delta Lake**, and **Bronze-Silver-Gold Medallion Architecture** to process large-scale Manufacturing IoT data.
- Built **incremental data processing** using **Delta MERGE, Watermarking**, and **Change Data Capture (CDC)** while optimizing **PySpark** workloads with **Partitioning, Caching, Broadcast Joins**, and **AQE**.
- Developed **metadata-driven ADF pipelines** with **dynamic parameterization** and integrated **Azure Key Vault** for secure credential and **Service Principal** management.
- Implemented **ACID Transactions, Schema Evolution, Time Travel, Source-to-Target Validation (STT)**, and **Data Quality Checks** to ensure reliable and accurate data processing.
- Collaborated with **Business Analysts, Solution Architects**, and **QA Teams** in an **Agile Scrum** environment to deliver scalable cloud-based data engineering solutions.

Data Engineer — SQL Server to Snowflake Enterprise Data Warehouse Modernization
Insurance Client

May 2022 – Oct 2024
Hyderabad, India.

- Developed an **Enterprise Data Warehouse Modernization** solution by migrating data from **SQL Server, Oracle**, and **Flat Files** into **Snowflake** using scalable **ELT** pipelines.
- Designed automated data ingestion using **Snowpipe, Streams, Tasks**, and **Snowpark**, supporting both full and **incremental data loads** across multiple source systems.
- Developed and optimized **Snowflake Stored Procedures, Views, UDFs, MERGE operations, CTEs**, and **Window Functions** to build high-performance analytical datasets.
- Configured **Virtual Warehouses, RBAC, Resource Monitors**, and managed deployments across **DEV, QA, UAT**, and **Production** environments.
- Optimized query performance using **Clustering Keys, Micro-Partition Pruning, Query Profile Analysis, Time Travel**, and **Zero Copy Cloning**, reducing ETL runtime from **80 minutes to 12 minutes**.
- Performed **Source-to-Target Validation (STT), Data Quality Checks**, and **Performance Tuning** to ensure accurate, secure, and reliable enterprise reporting.

Projects

SnowConvert AI – SQL Server to Snowflake Migration Framework

- Developed an automated **SQL Server to Snowflake Migration Framework** using **SnowConvert AI** to assess migration complexity, convert **Tables, Views, Stored Procedures**, and **Functions**, while generating detailed migration assessment reports.
- Built an end-to-end validation framework using **Snowflake, Python**, and **SQL** to perform **Source-to-Target Validation (STT), Data Quality Checks**, and **Performance Benchmarking** for migrated workloads.

Enterprise Sales Analytics Platform

- Designed an end-to-end **Microsoft Fabric** analytics platform leveraging **OneLake, Fabric Data Factory, Fabric Lakehouse**, and **Fabric Data Warehouse** to centralize enterprise sales data.
- Implemented **PySpark-based ETL/ELT** pipelines using **Fabric Notebooks, Delta Lake, Bronze-Silver-Gold Medallion Architecture**, and **Incremental Loading** to build curated analytical datasets.
- Developed interactive **Power BI** dashboards and optimized **Fabric Data Warehouse** performance through **Data Modelling, Star Schema**, and **Query Optimization**, enabling real-time business reporting.